

March 1, 2021

DNS



- Today
 - Paper 2 Info
 - DNS, the good, the bad and the ugly
 - Paper Presentation
- Assignments
 - Project
 - Outline: Due Monday, Mar 29

Paper 2

- Need selection of dates (1st and 2nd choices) by next week
- Choice of paper must be approved by Dr C at least two weeks prior to the presentation
 - Preference is for current (no more than 2-3 years) academic journal or conference article
 - Commercial whitepapers are acceptable unless they're pure marketing with no technical content
 - If possible, paper should support your project choice.
- Choice of paper will be announced to the class at least one week prior to the presentation

Domain Name System

- Transaction Structure
- Message Flow
- Proxies and Caches

DNS

- DNS is distributed
 - Organized as a tree, with the **root nameservers** at the top
 - Each **top-level domain (TLD)** (e.g., .com, .edu, .gov, .uk) served by a separate root nameserver
- Authoritative Name Servers responsible for their domains
- Domain information stored as a **zone record**

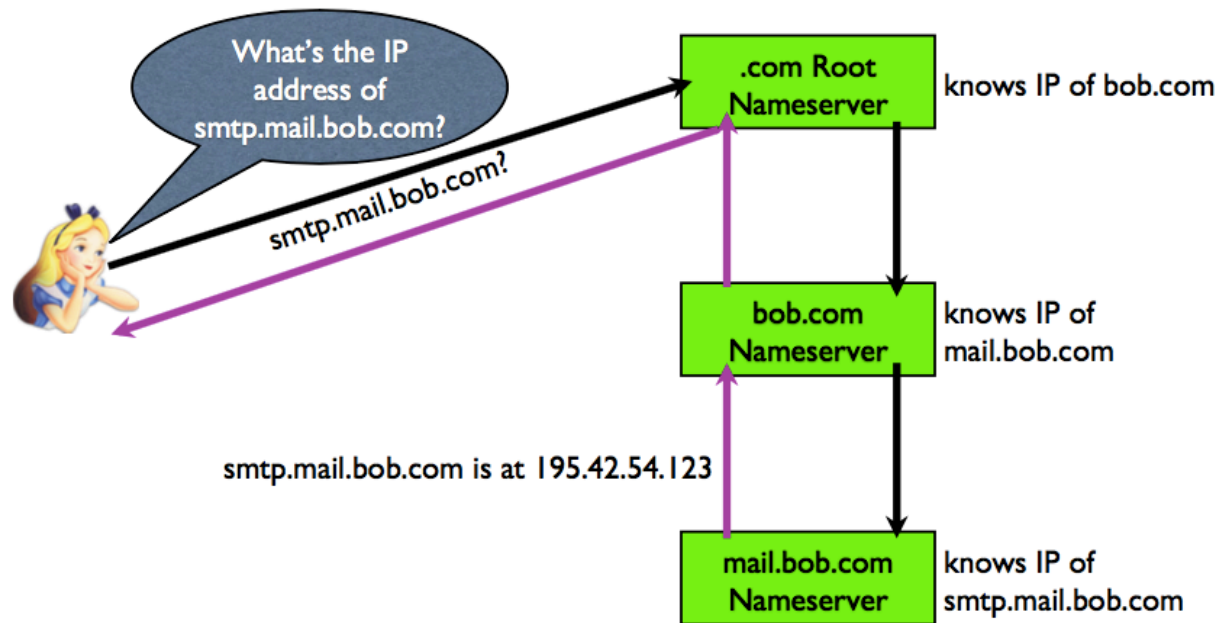
Name Servers

- **AuthoritativeNameServer**: gives authoritative results for hostnames that have been configured
- Domains are registered with a **domain name registrar** (e.g., GoDaddy)
- Each domain must have one primary and at least one secondary name servers
- For reliability in case of failure

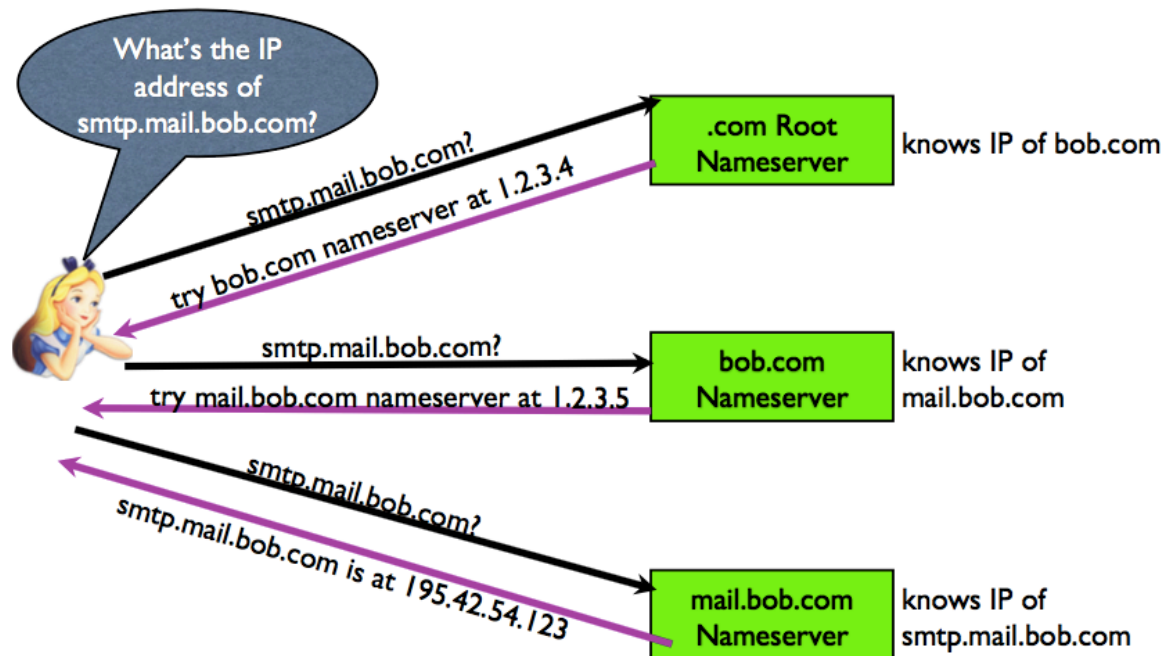
TLDs

- Name servers pre-loaded with IP addresses of TLD name servers

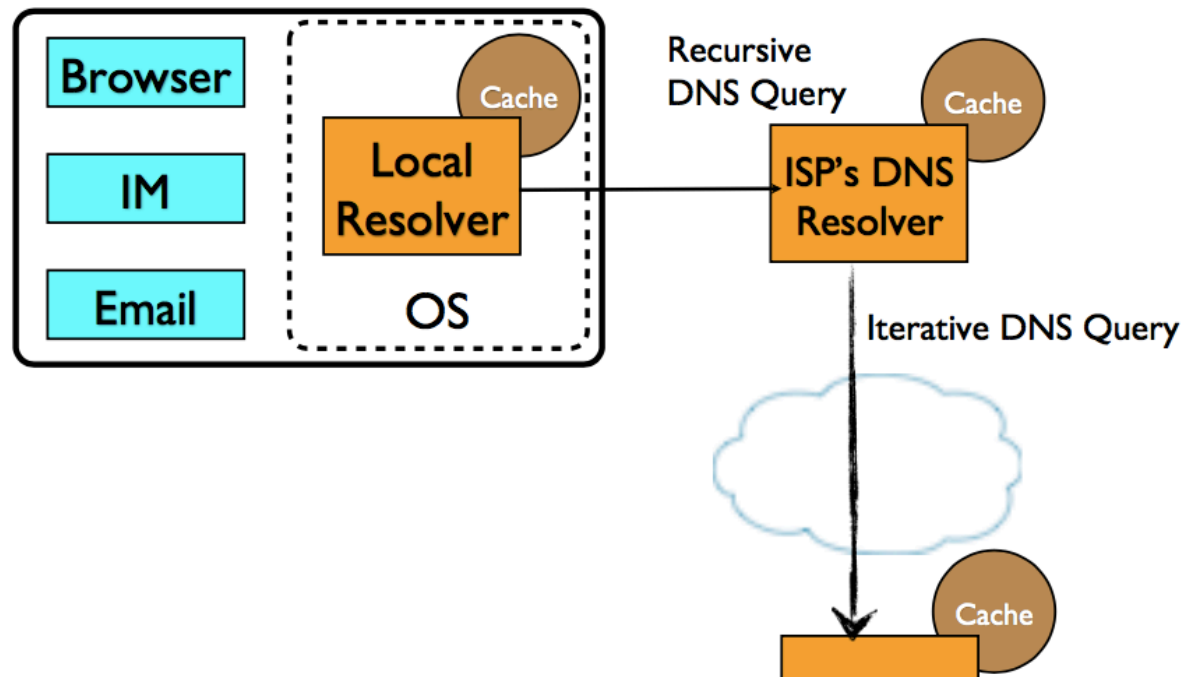
Naïve Recursive Query



Naïve Iterative Query



DNS in the real world



DNS Vulnerabilities

- DNS requests and responses are not authenticated
 - Yet many applications trust DNS resolutions
 - ... or, more accurately, they don't consider the threat at all
 - Spoofing of DNS is very dangerous -- **WHY?**
- Caching doesn't help:
 - DNS relies heavily on caching for efficiency, enabling **cache pollution** attacks
 - Once something is wrong, it can remain that way in caches for a long time
 - Data may be corrupted before it gets to authoritative server

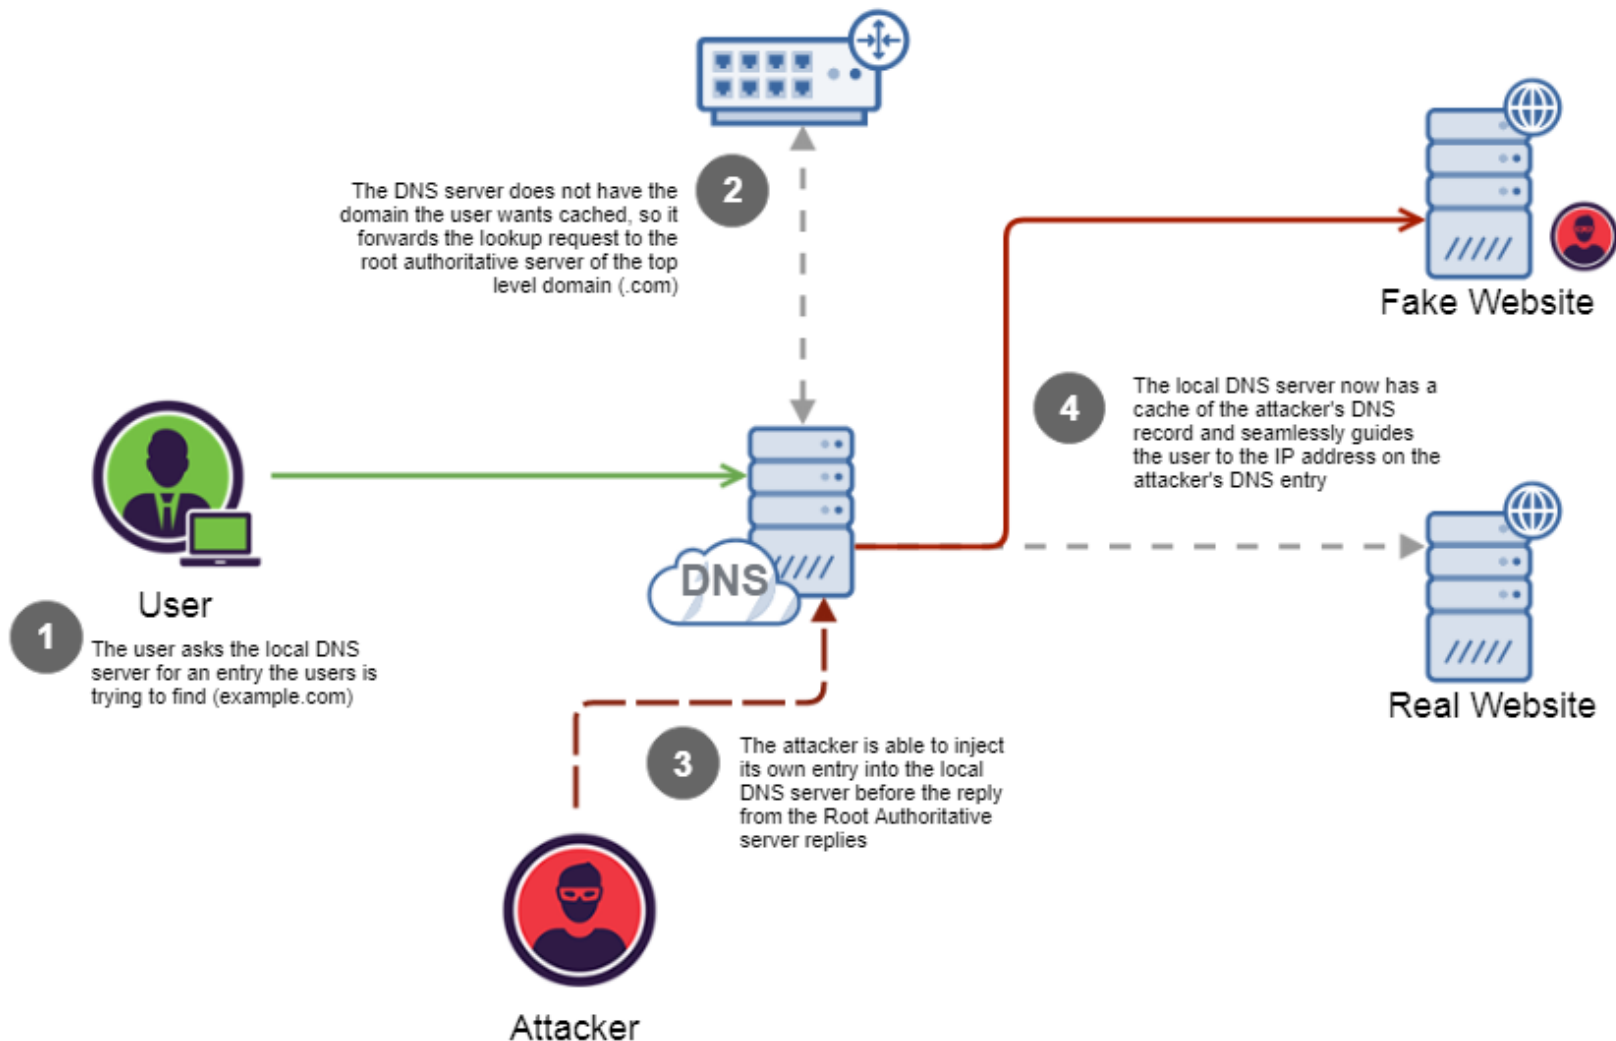
A Cache Poisoning Attack

- All DNS requests have a unique query ID
- The nameserver/resolver uses this information to match up requests and responses -- this is useful since DNS uses UDP
- If an adversary can guess the query ID, then it can forge the responses and pollute the DNS cache
 - 16-bit query IDs (only $2^{16}=65536$ possible query IDs)
 - Some servers increment IDs (or use some other predictable algo)
 - gethostbyname returns as soon as it gets a response, so first one in wins!!!
- Note: If you can observe the traffic going to a name server, you can pretty much arbitrarily Own the Internet for the clients it serves

A Cache Poisoning Attack

- A simple (and extremely effective) attack:
 1. Wait for Alice to send DNS request to nameserver
 2. Intercept request
 3. Quickly insert a fake response
- If attacker is faster and/or closer to Alice than the DNS server, then the attack is successful
- Advantage attacker: unlike the name server, the attacker doesn't have to do any actual resolving

Root Authoritative DNS



Attack Limitations

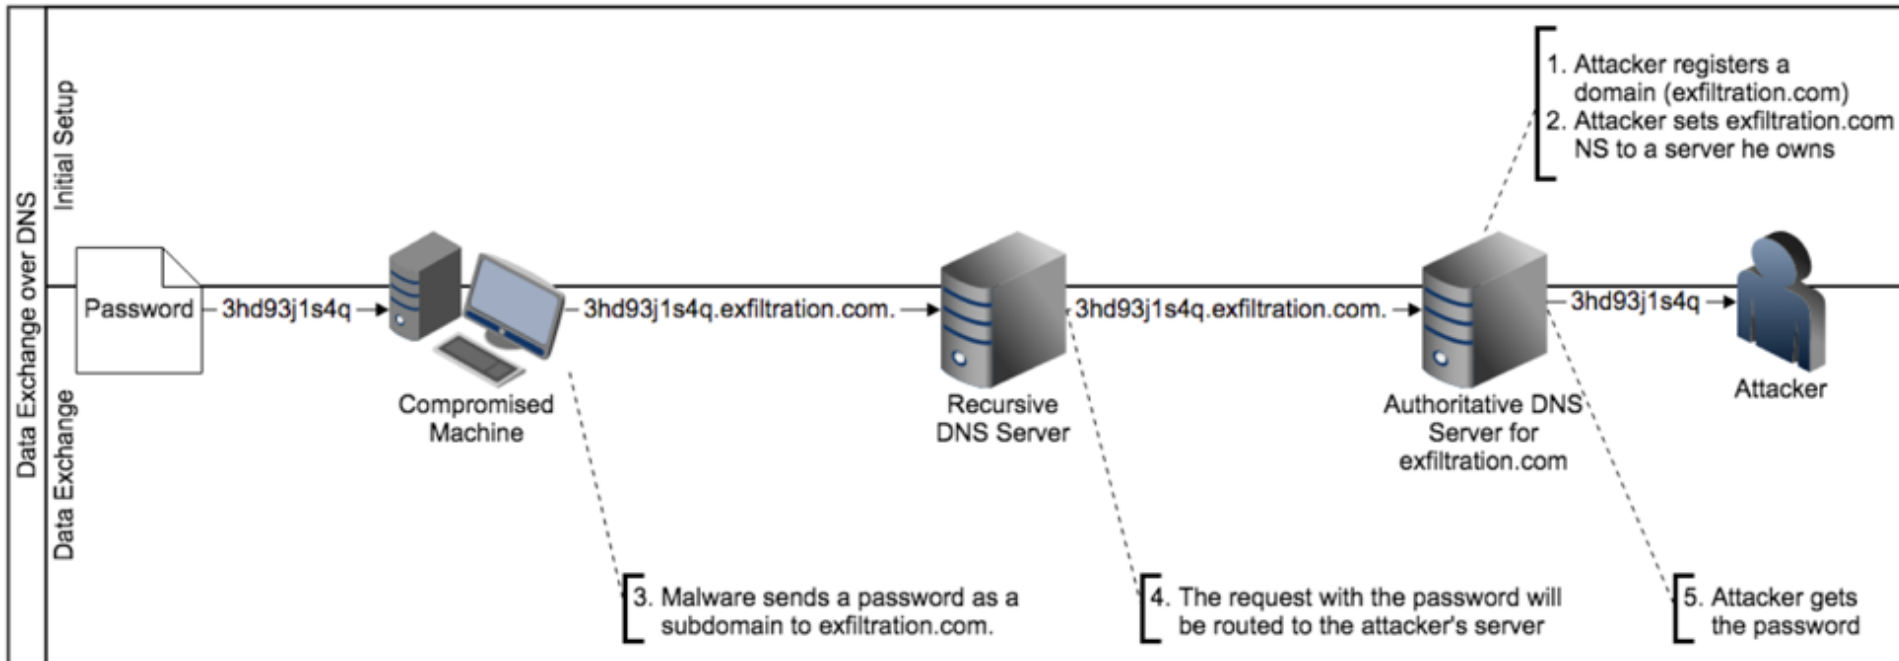
- Victim hostname cannot already be in the cache
- Randomizing the QueryID makes the race condition much harder to exploit
(2^{16} possible QueryIDs)

Kaminsky Attack

- Hijacks the entire name server of victim host • Basic idea
- Choose a random hostname in the domain (guaranteed not to be cached)
- Try to beat real name server response (guessing the QueryID)
- Forged response specifies an update for the name server IP address (to attacker)
- Repeat until successful
- All future DNS queries for the victim domain now directed to the attacker's DNS server (until TTL expires)

DNS for Exfiltration

- Assuming the Kaminsky attack has succeeded
 - Or (think about the structure of DNS)
- What can we do with it?



Mitigations?

- The QueryID is 16 bits.
- Increasing the size would break the Internet
- What else can we randomize?
 - Source port address

DNS Going Forward

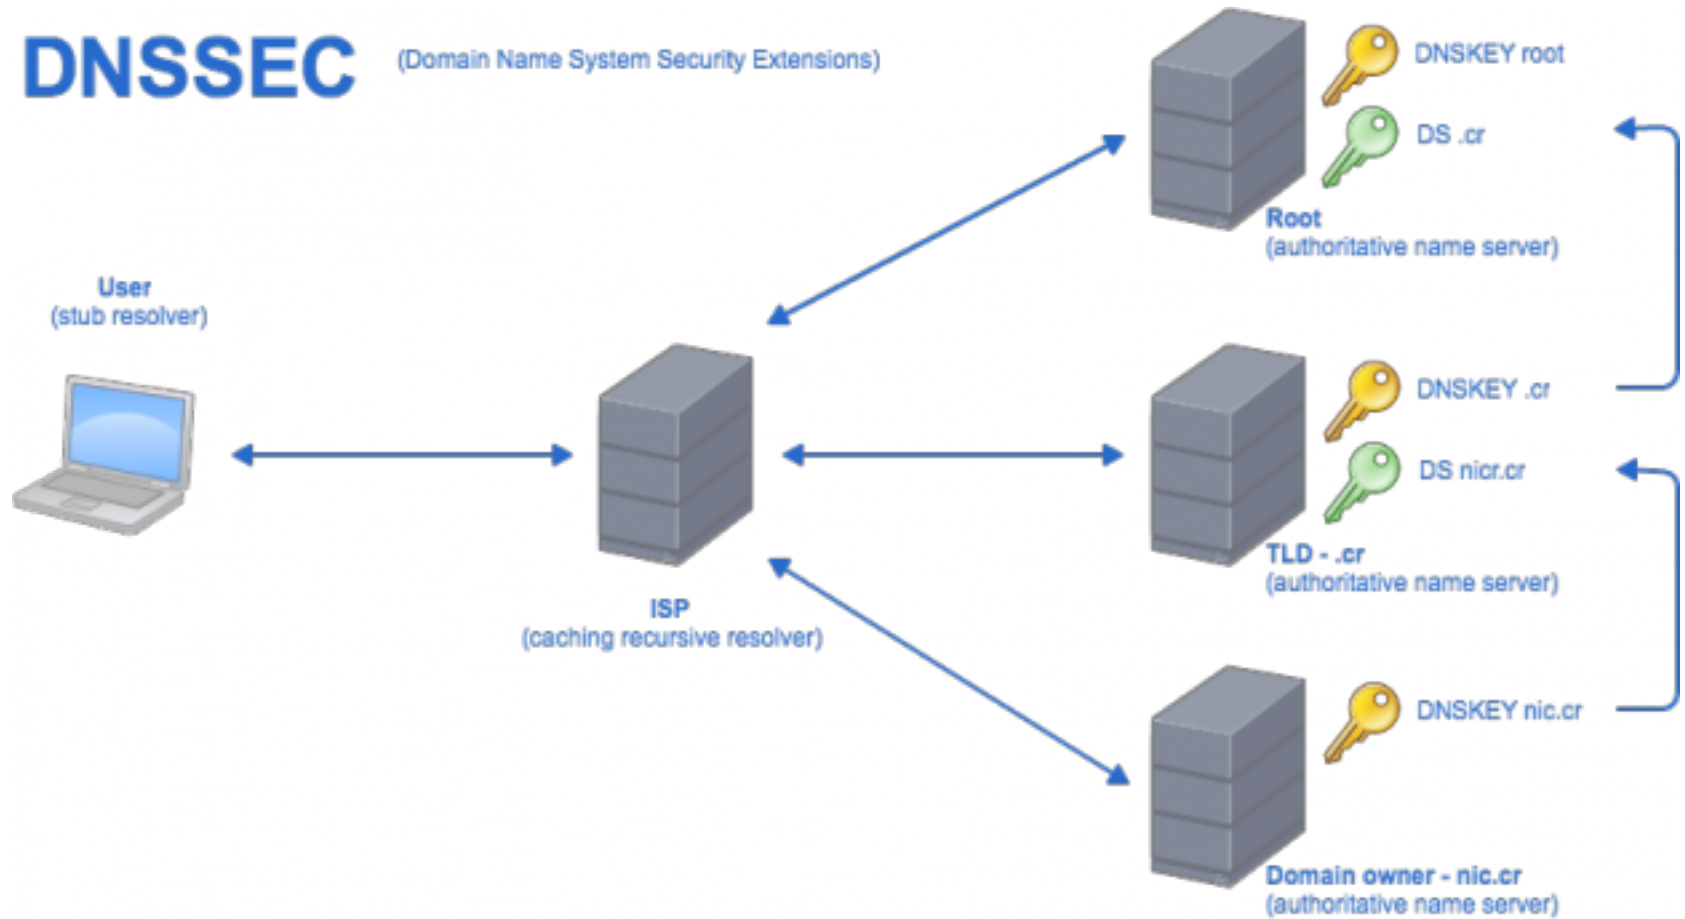
DNSSec

A standards-based (IETF) solution to security in DNS

- Prevents data spoofing and corruption
- Authentication (verifiable DNS) using public key infrastructure
- Authenticates:
 - Communication between servers
 - DNS data
 - Content
 - Existence
 - non-existence
 - Public keys


DNSSEC

(Domain Name System Security Extensions)



DNSSEC Mechanisms

- Each domain signs their “zone” with a private key
- Public keys published via DNS
- Zones signed by parent zones
- Ideally, you only need a self-signed root, and follow keys down the hierarchy



DNSSEC Challenges

- Incremental deployability
 - Everyone has DNS, can't assume a flag day
- Resource imbalances
 - Some devices can't afford real authentication
- Performance
 - Certificate process
- Cultural
 - Who gets to control the root keys? (US, China, CofC?)
 - Most people don't have any strong reason to have secure DNS (\$\$\$ not justified in most environments)
 - Lots of transitive trust assumptions
- Take away: DNSSEC will be deployed, but it is unclear whether it will be used appropriately/widely